

# DSA-2000



Real-time aperture synthesis  
imaging with the DSA-2000 radio camera

2024 Legion Retreat

Martin Pokorny, Caltech

[www.deepsynoptic.org](http://www.deepsynoptic.org)

Caltech



Schmidt Sciences

# What is DSA-2000?

## A world-leading radio survey telescope and multi-messenger discovery engine.

- ~2000 x 5m dishes (19 x 15 km)
- Spring Valley, Nevada
- Frequency: 0.7 - 2 GHz band
- Spatial resolution: 3.3 arcseconds
- *Highly optimized for surveys*

Currently in final design phase

Construction start: 2026

First light: 2028

Key surveys: 2028 – 2033



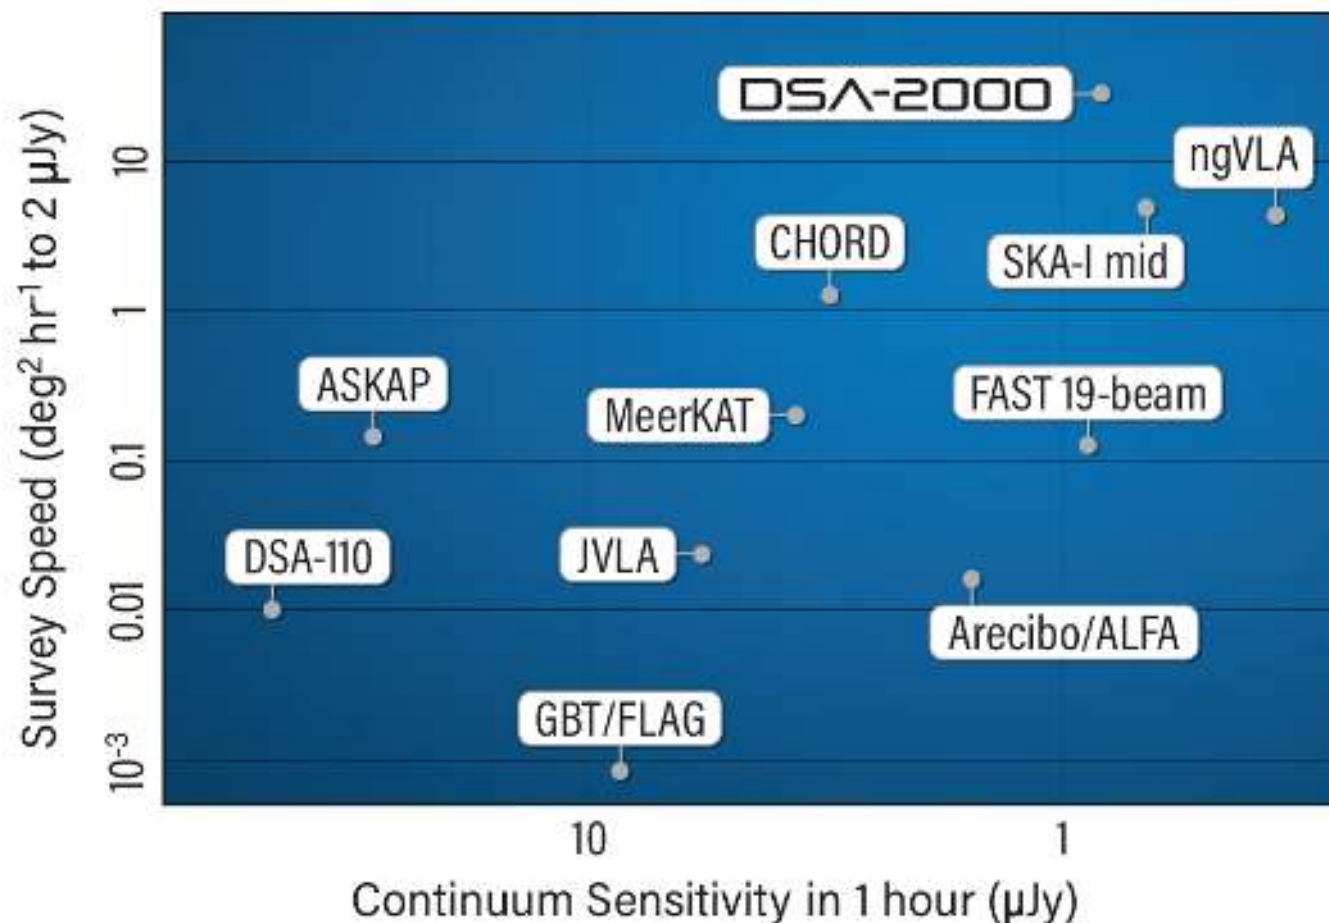


## Unparalleled survey speed

31,000 deg<sup>2</sup> to 500 nJy  
 >1 billion radio sources (IQUV)  
 ~few million galaxies in HI  
 ~10<sup>5</sup> FRBs and pulsars  
 ~10<sup>6</sup> “slow” transients

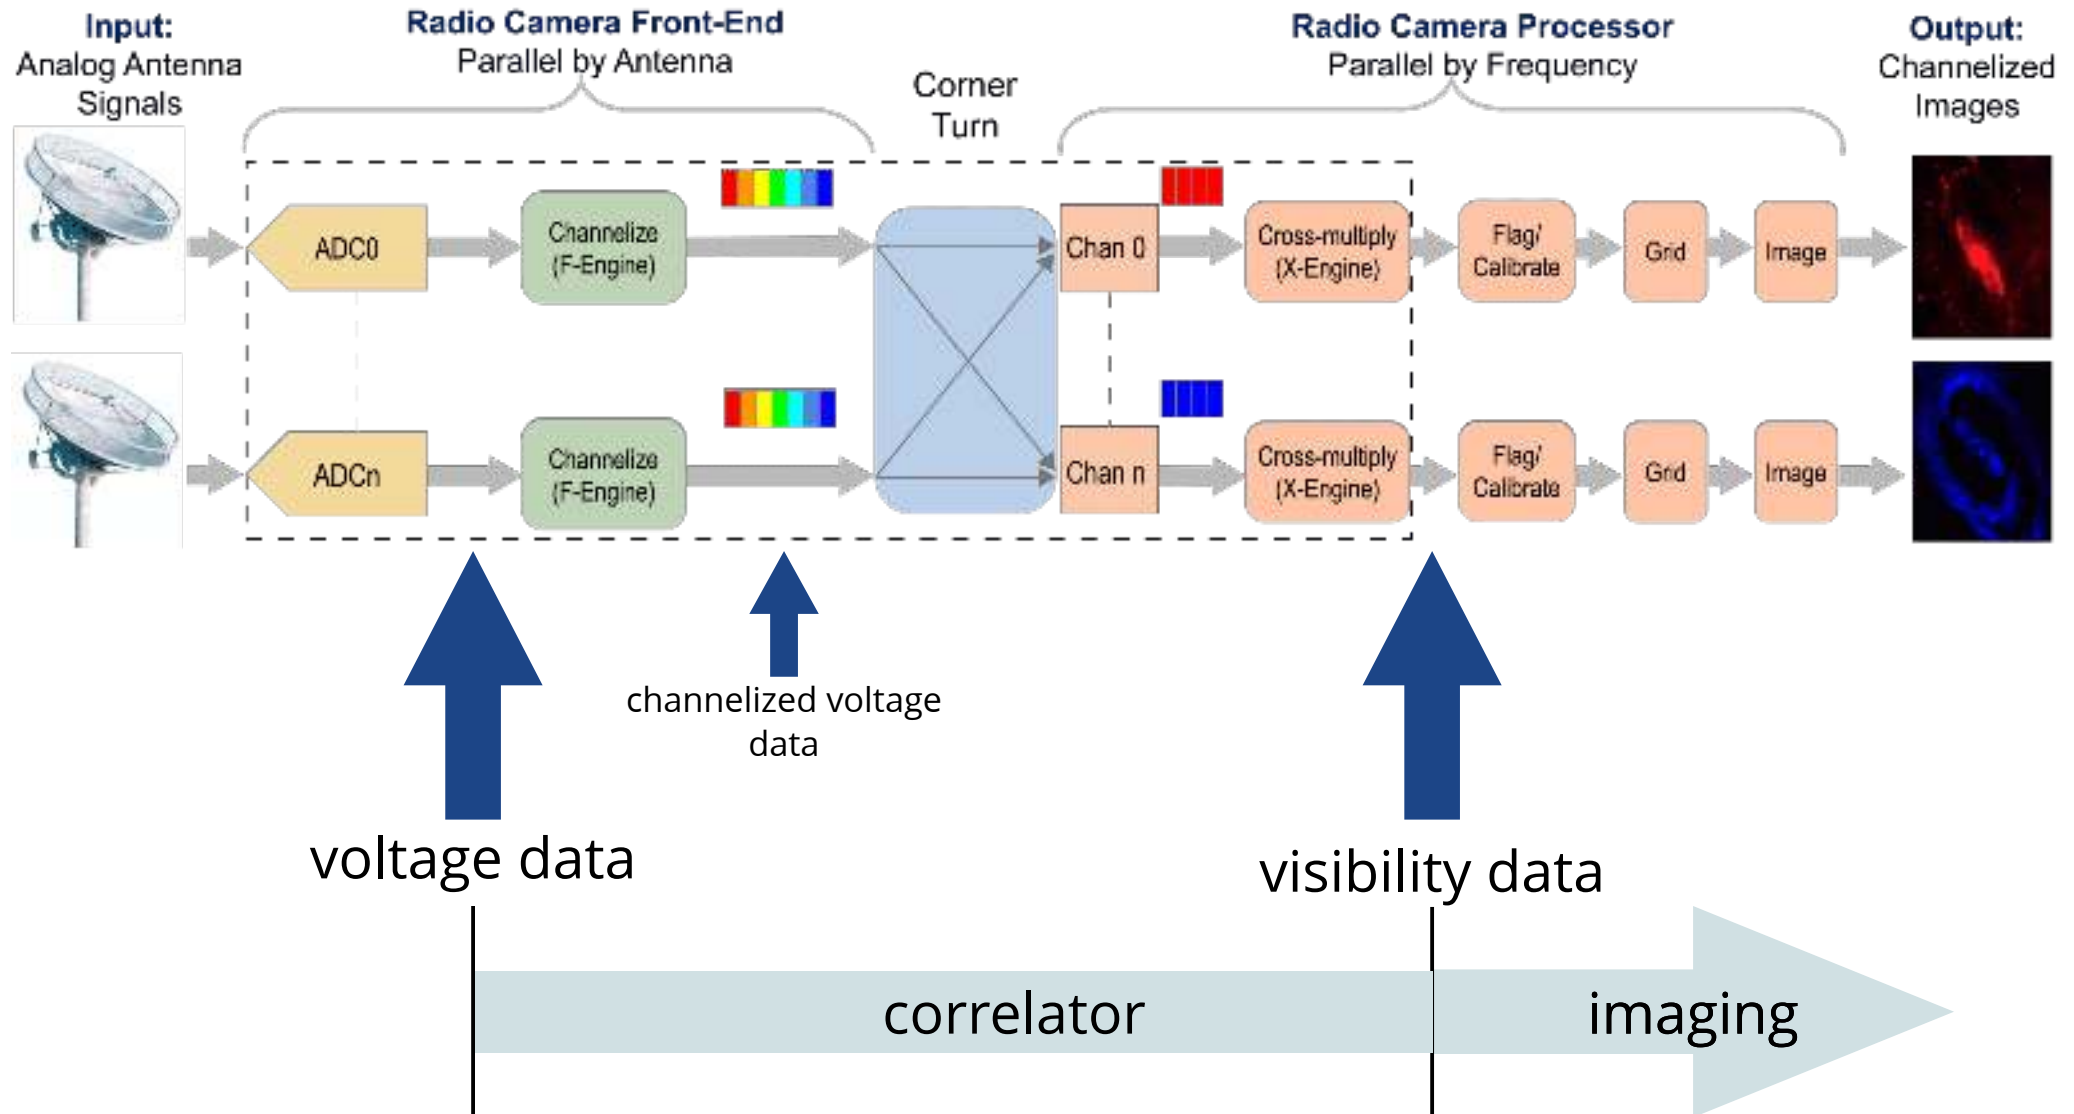
Enabled by two key technologies:

- “Radio camera” digital back-end
- Cryo-free antenna/receiver




# What is a radio camera?

## Radio camera schematic



## Radio camera motivation


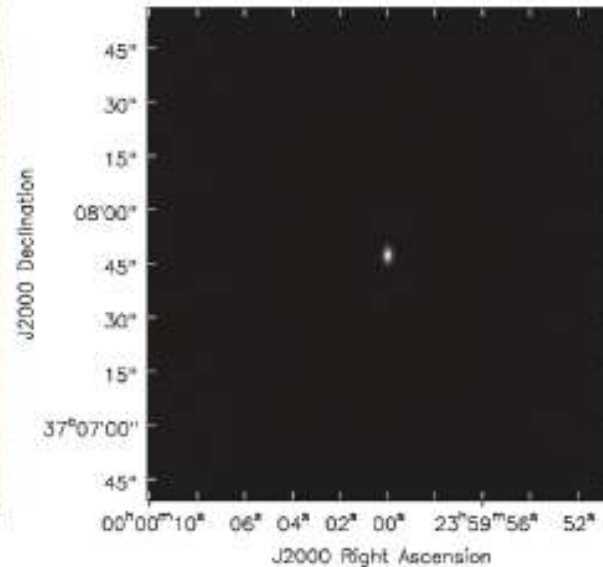


Radio telescope correlator output



### DSA-2000

- ~0.8 Zettabytes/yr of voltage data
- 20 Exabytes/yr of visibility data
- \$100 million/yr in storage costs to save visibility data

## Radio camera imaging



- No need for visibility-based deconvolution
- Enables a deterministic stream processing pipeline that creates images



## Radio camera processing characteristics

- High degree of data parallelism, by frequency channel slices (sub-bands)
  - no cross-sub-band communication, except, maybe...
- Calibration improves with greater bandwidth; may need to span sub-bands.
- RCP input (channelized voltages): 43Tb/s
- RCP output (images): 23 GB/s

Data	Type	Rate (per input channel)
voltages	4+4 bit complex integer	4 Gb/s
raw visibilities	32+32 bit complex integer	43 MB/s
normalized visibilities	complex single precision fp	43 MB/s
gridded visibilities	complex single precision fp	107 MB/s
grid pixel (updates)	complex double precision fp	~43 GB/s (20x20 CF kernel)

## Radio camera processor deployment

- Goal: 4 input channels per GPU
- All 20 channels in each sub-band are processed by a single node
- 10 GPUs per node: 2 sub-bands per node
  - 3800 GPUs
- baseline GPU: NVIDIA RTX 4000 Ada generation
- 200 GbE network interface (for data input)
- Image output to node-local file system for post-processing

# Imaging pipeline software

*symcam*

(SYnthesis iMaging CAMera)

# Imaging pipeline inputs


- Data (channelized voltages)
  - for every channel, data arrive from all antennas
  - one input data "stream" for every N (4) contiguous channels
  - each antenna multicasts voltage data *via* UDP on all streams
    - raw data: 7.5  $\mu$ sec sample time
    - packetized data: several sample times (maximum possible for MTU)
  - metadata: timestamp, stream identification (channels), flag values (bad data)
- Events
  - control messages (*e.g.*, start/stop imaging, run bandpass calibration, start/stop diagnostic output)
  - system status metadata (*e.g.*, antenna beam patterns, calibration terms)
  - metadata: activation/validity timestamp




# From inputs to regions

- Both data and events are received and buffered by dedicated threads
- Legion tasks launched from the main loop fill regions associated with timestamps with data or events from service thread buffers
- Real-time conditions
  - data timeout: derived from timestamps of received packets
  - event deadline: derived from completion of tasks that fill data buffers
  - no explicit system clock dependency except when no data are received


## Main loop: input management



## Main loop: input management




## Main loop: input management







## Main loop: input management



# Main loop: input management

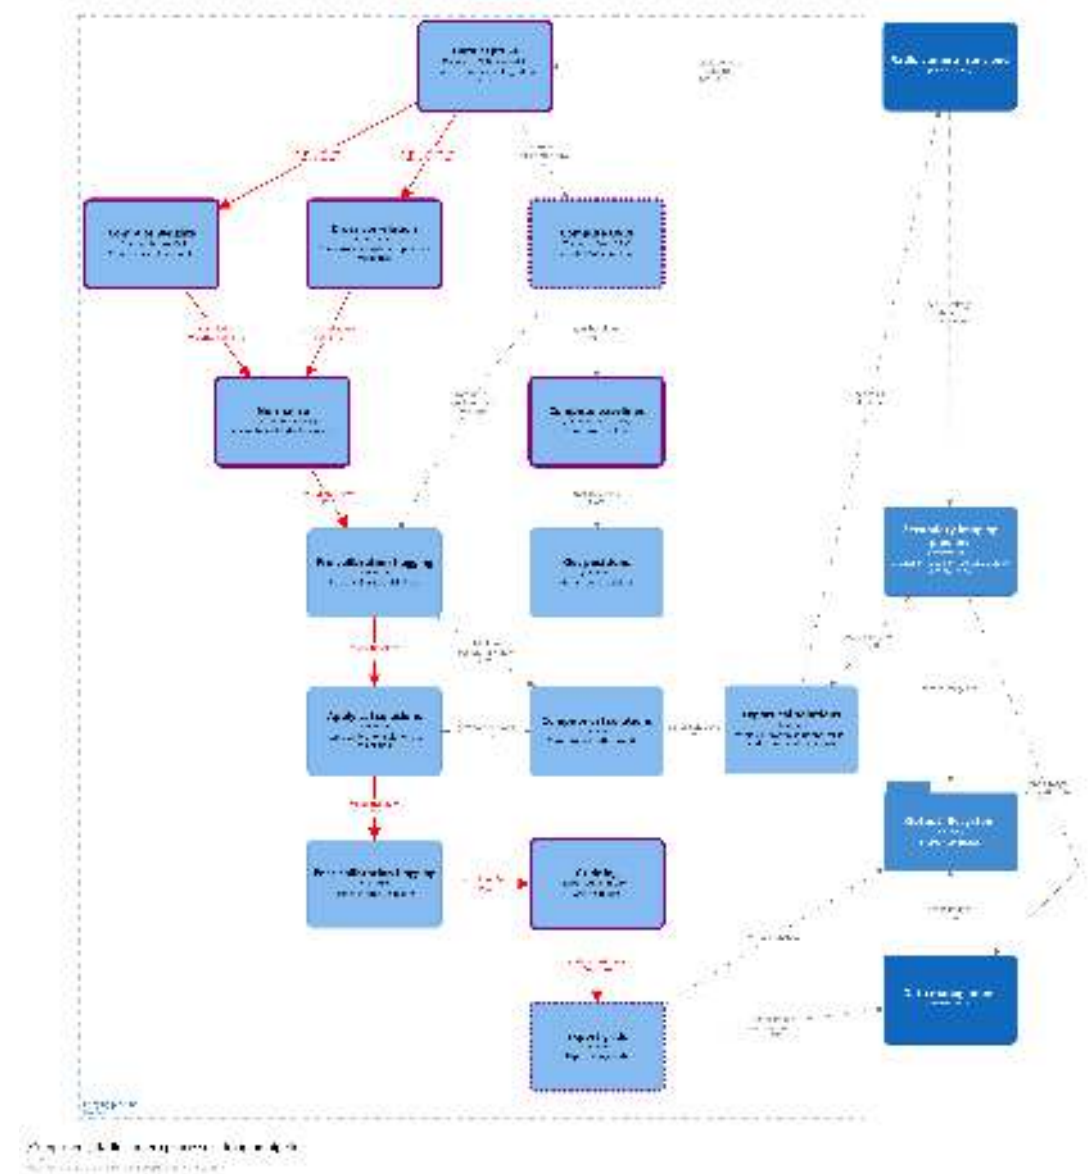


## Main loop: input management




# Pipeline states and predication

- a pipeline state is associated with every data timestamp
- states are determined by events
- states determine processing for data with associated timestamps
- pipeline processing is a series of stages, implemented as tasks
- states provide predicates for tasks, effectively making each stage conditional on the pipeline state





# Channel data regions & partitions



# Current status and next steps

# Status

- Events
  - create/destroy subarray
  - shutdown
  - start/stop imaging
  - start/stop flagging
  - start/stop solve calibration
  - start/stop apply calibration (partial)
  - start/stop record data
  - start/stop record events
  - import/export fftw wisdom

# Status

- Events
  - create/destroy subarray
  - shutdown
  - start/stop imaging
  - start/stop flagging
  - start/stop solve calibration
  - start/stop apply calibration (partial)
  - start/stop record data
  - start/stop record events
  - import/export fftw wisdom
- Stages
  - data capture
  - cross-correlation
  - data weight computation (flag correlation)
  - normalization
  - gridding
  - Fourier transform
  - image export (partial)



# Status

- Events
  - create/destroy subarray
  - shutdown
  - start/stop imaging
  - start/stop flagging
  - start/stop solve calibration
  - start/stop apply calibration (partial)
  - start/stop record data
  - start/stop record events
  - import/export fftw wisdom
- Stages
  - data capture
  - cross-correlation
  - data weight computation (flag correlation)
  - normalization
  - gridding
  - Fourier transform
  - image export (partial)
- Incomplete
  - data flagging (RFI)
  - solve calibration
    - modal (delay, bandpass, polarization)
    - continuous (gain)
  - apply calibration
  - bright source subtraction
  - export calibration solutions
  - import calibration solutions

# Integration with calibration

## Current situation

- calibration algorithms developed by domain expert
- implemented using Python
  - numpy
  - JAX
  - Ray
- use inter-process comm with a Ray actor for computing calibration solutions
  - ugly, but may work if averaging is done on Legion side

# Integration with calibration

## Current situation

- calibration algorithms developed by domain expert
- implemented using Python
  - numpy
  - JAX
  - Ray
- use inter-process comm with a Ray actor for computing calibration solutions
  - ugly, but may work if averaging is done on Legion side

## Future design

- desirable, but perhaps unnecessary
- Legate
- tasks implemented using XLA

## Main loop refactor

- main challenge of main loop is to run -- not just launch -- data capture tasks ahead of time to avoid losing input data
- restrict main loop to issuing data and event capture tasks, and a single, high-level "input data processing" task, which runs after data and required events have arrived
- all states are known when input data processing task runs -- no predication needed